

# investigate-a-dataset-template.ip

January 3, 2018

## 1 Project: Investigate a Dataset (Replace this with something more specific!)

### 1.1 Table of Contents

#### Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

## Introduction

The selected dataset is the TMDb dataset which contains the data on movies as well as ratings

Exploration of the following trends will be done: > Runtime of movies over the years > Popularity of movies over the years > Revenue of movies over the years

Associations of various factors to be seen are: > Revenue vs Popularity > Runtime vs Popularity > Runtime vs Revenue

```
In [77]: # Use this cell to set up import statements for all of the packages that you
        #      plan to use.
```

```
# Remember to include a 'magic word' so that your visualizations are plotted
# inline with the notebook. See this page for more:
# http://ipython.readthedocs.io/en/stable/interactive/magics.html
import pandas as pd
from datetime import datetime
import matplotlib.pyplot as plt
import seaborn as sns
% matplotlib inline
```

## 2

### 2.1 Data Wrangling

**Tip:** In this section of the report, you will load in the data, check for cleanliness, and then trim and clean your dataset for analysis. Make sure that you document your steps carefully and justify your cleaning decisions.

### 2.1.1 General Properties

```
> Rows: 4813
> Columns: 18
```

Null values are present in variables which are as follows: imdb\_id cast homepage director tagline keywords overview genres production\_companies

Steps for cleaning the data:

- 1.Data seems to be relatively clean
- 2.one duplicate record is found and is deleted
- 3.Of the different variables, imdb\_id is significant to identify the records uniquely. Only 10 of 10865 are missing. Hence, we only drop the records with this value as missing.
- 4.Some other values with missing values are significant so they are retained

```
In [6]: pwd
```

```
Out[6]: 'C:\\Users\\hi\\Desktop\\dand pro 3'
```

```
In [7]: # Load your data and print out a few lines. Perform operations to inspect data
#       types and look for instances of missing or possibly errant data.
df = pd.read_csv('tmdb_movies.csv',encoding='ISO-8859-1')
df.head(3)
```

```
Out[7]:
```

	id	imdb_id	popularity	budget	revenue	original_title
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent

```
cast \
0 Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
1 Tom Hardy|Charlize Theron|Hugh Keays-Byrne|Nic...
2 Shailene Woodley|Theo James|Kate Winslet|Ansel...
```

```
homepage director \
0 http://www.jurassicworld.com/ Colin Trevorrow
1 http://www.madmaxmovie.com/ George Miller
2 http://www.thedivergentseries.movie/#insurgent Robert Schwentke
```

```
tagline ... \
0 The park is open. ...
1 What a Lovely Day. ...
2 One Choice Can Destroy You ...
```

```
overview runtime \
0 Twenty-two years after the events of Jurassic ... 124
1 An apocalyptic story set in the furthest reach... 120
2 Beatrice Prior must confront her inner demons ... 119
```

```
genres \
```

```

0 Action|Adventure|Science Fiction|Thriller
1 Action|Adventure|Science Fiction|Thriller
2      Adventure|Science Fiction|Thriller

```

```

                                production_companies release_date vote_count \
0 Universal Studios|Amblin Entertainment|Legenda...      6/9/15      5562
1 Village Roadshow Pictures|Kennedy Miller Produ...      5/13/15      6185
2 Summit Entertainment|Mandeville Films|Red Wago...      3/18/15      2480

```

```

      vote_average  release_year  budget_adj  revenue_adj
0              6.5           2015  1.379999e+08  1.392446e+09
1              7.1           2015  1.379999e+08  3.481613e+08
2              6.3           2015  1.012000e+08  2.716190e+08

```

```
[3 rows x 21 columns]
```

```
In [8]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64
vote_average      10866 non-null float64
release_year      10866 non-null int64
budget_adj        10866 non-null float64
revenue_adj       10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

```

### 3 Finding and dealing with Duplicates

Here, we check for duplicates. These would be redundant records for our investigation and if present, can be dropped from the dataset

```
In [9]: # checking number of duplicate records
        sum(df.duplicated())
```

```
Out[9]: 1
```

```
In [10]: # drop rows that are duplicates
          # save in dataset itself, so that data is devoid of duplicates
          df.drop_duplicates(inplace=True)
```

```
In [11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 21 columns):
id                10865 non-null int64
imdb_id           10855 non-null object
popularity        10865 non-null float64
budget            10865 non-null int64
revenue           10865 non-null int64
original_title    10865 non-null object
cast              10789 non-null object
homepage          2936 non-null object
director          10821 non-null object
tagline           8041 non-null object
keywords          9372 non-null object
overview          10861 non-null object
runtime           10865 non-null int64
genres            10842 non-null object
production_companies 9835 non-null object
release_date      10865 non-null object
vote_count        10865 non-null int64
vote_average      10865 non-null float64
release_year      10865 non-null int64
budget_adj        10865 non-null float64
revenue_adj       10865 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.8+ MB
```

### 4 NULL Values:

This step is to inspect if data is sufficient for initial exploration and research. Depending on the results, we also will have to decide on how best to tackle missing values, based on their relative proportion to complete data.

```
In [12]: # check if any columns contain null values
df.isnull().sum()
```

```
Out[12]: id                0
         imdb_id           10
         popularity        0
         budget            0
         revenue           0
         original_title     0
         cast              76
         homepage          7929
         director           44
         tagline            2824
         keywords           1493
         overview           4
         runtime            0
         genres             23
         production_companies 1030
         release_date        0
         vote_count          0
         vote_average        0
         release_year        0
         budget_adj          0
         revenue_adj         0
         dtype: int64
```

#### 4.0.1 Data Cleaning (Dropping some rows)

1. After viewing the TMDb dataset and deleting duplicate records, we perform cleaning steps in this part

2. As we saw from the previous step that we have null values in certain columns. Of these, `imdb_id` is of relevance to us. The number of records with missing `imdb_id` values are only 10 out of 10865. Hence, we are removing rows with null values only in `imdb_id` column using `dropna` method.

3. Remaining columns containing missing values are not factors of interest in this initial exploration. The questions we have posed in our EDA consider factors having no missing values. So, we leave the remaining missing values.

```
In [13]: #dropping rows containing missing values in imdb_id column
df.dropna(subset=['imdb_id'], inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10855 entries, 0 to 10865
Data columns (total 21 columns):
id                10855 non-null int64
imdb_id           10855 non-null object
popularity        10855 non-null float64
budget            10855 non-null int64
```

```

revenue          10855 non-null int64
original_title   10855 non-null object
cast             10779 non-null object
homepage         2934 non-null object
director         10815 non-null object
tagline          8038 non-null object
keywords         9368 non-null object
overview         10852 non-null object
runtime          10855 non-null int64
genres           10834 non-null object
production_companies 9830 non-null object
release_date     10855 non-null object
vote_count       10855 non-null int64
vote_average     10855 non-null float64
release_year     10855 non-null int64
budget_adj       10855 non-null float64
revenue_adj      10855 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.8+ MB

```

```

In [14]: #checking dataset
df.head()

```

```

Out[14]:
   id  imdb_id  popularity  budget  revenue \
0  135397  tt0369610   32.985763  150000000  1513528810
1   76341  tt1392190   28.419936  150000000   378436354
2  262500  tt2908446   13.112507  110000000   295238201
3  140607  tt2488496   11.173104  200000000  2068178225
4  168259  tt2820852    9.335014  190000000  1506249360

```

```

   original_title \
0      Jurassic World
1    Mad Max: Fury Road
2      Insurgent
3  Star Wars: The Force Awakens
4      Furious 7

```

```

   cast \
0  Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
1  Tom Hardy|Charlize Theron|Hugh Keays-Byrne|Nic...
2  Shailene Woodley|Theo James|Kate Winslet|Ansel...
3  Harrison Ford|Mark Hamill|Carrie Fisher|Adam D...
4  Vin Diesel|Paul Walker|Jason Statham|Michelle ...

```

```

   homepage          director \
0  http://www.jurassicworld.com/  Colin Trevorrow
1  http://www.madmaxmovie.com/    George Miller

```

```

2      http://www.thedivergentseries.movie/#insurgent Robert Schwentke
3      http://www.starwars.com/films/star-wars-episod... J.J. Abrams
4      http://www.furious7.com/ James Wan

```

```

tagline ... \
0      The park is open. ...
1      What a Lovely Day. ...
2      One Choice Can Destroy You ...
3      Every generation has a story. ...
4      Vengeance Hits Home ...

```

```

overview runtime \
0      Twenty-two years after the events of Jurassic ... 124
1      An apocalyptic story set in the furthest reach... 120
2      Beatrice Prior must confront her inner demons ... 119
3      Thirty years after defeating the Galactic Empi... 136
4      Deckard Shaw seeks revenge against Dominic Tor... 137

```

```

genres \
0      Action|Adventure|Science Fiction|Thriller
1      Action|Adventure|Science Fiction|Thriller
2      Adventure|Science Fiction|Thriller
3      Action|Adventure|Science Fiction|Fantasy
4      Action|Crime|Thriller

```

```

production_companies release_date vote_count \
0      Universal Studios|Amblin Entertainment|Legenda... 6/9/15 5562
1      Village Roadshow Pictures|Kennedy Miller Produ... 5/13/15 6185
2      Summit Entertainment|Mandeville Films|Red Wago... 3/18/15 2480
3      Lucasfilm|Truenorth Productions|Bad Robot 12/15/15 5292
4      Universal Pictures|Original Film|Media Rights ... 4/1/15 2947

```

```

vote_average release_year budget_adj revenue_adj
0      6.5      2015 1.379999e+08 1.392446e+09
1      7.1      2015 1.379999e+08 3.481613e+08
2      6.3      2015 1.012000e+08 2.716190e+08
3      7.5      2015 1.839999e+08 1.902723e+09
4      7.3      2015 1.747999e+08 1.385749e+09

```

[5 rows x 21 columns]

```

In [22]: # save new cleaned dataset.
# we will use this dataset in next sections
df.to_csv('data_imdb.csv', index=False)

```

## Exploratory Data Analysis

#### 4.0.2 Research Question 1:

what are the trends of runtimes, popularity and revenues over time ?

1. For this analysis, we first group our data based on years, using the variable 'release\_year'

2. After that we use the built in histogram function to visualize our answer.


```
In [27]: # Use this, and more code cells, to explore your data. Don't forget to add
#         Markdown cells to document your observations and findings.
df_imdb = pd.read_csv('data_imdb.csv', encoding='ISO-8859-1')
df_imdb.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10855 entries, 0 to 10854
Data columns (total 21 columns):
id                10855 non-null int64
imdb_id           10855 non-null object
popularity        10855 non-null float64
budget           10855 non-null int64
revenue          10855 non-null int64
original_title    10855 non-null object
cast             10779 non-null object
homepage         2934 non-null object
director         10815 non-null object
tagline          8038 non-null object
keywords         9368 non-null object
overview         10852 non-null object
runtime          10855 non-null int64
genres           10834 non-null object
production_companies 9830 non-null object
release_date     10855 non-null object
vote_count       10855 non-null int64
vote_average     10855 non-null float64
release_year     10855 non-null int64
budget_adj       10855 non-null float64
revenue_adj      10855 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

```
In [30]: #group the input data by release_year variable and compute mean values for the numerical
df_new = df_imdb.groupby('release_year').mean()
```


```
In [31]: #Initial exploration to see possible trends visually
df_new.hist(figsize=(10,10));
```





In [84]: *# summary of the dataset using seaborn which is used for visualization*  
`sns.pairplot(df_new,size=2,hue='popularity')`

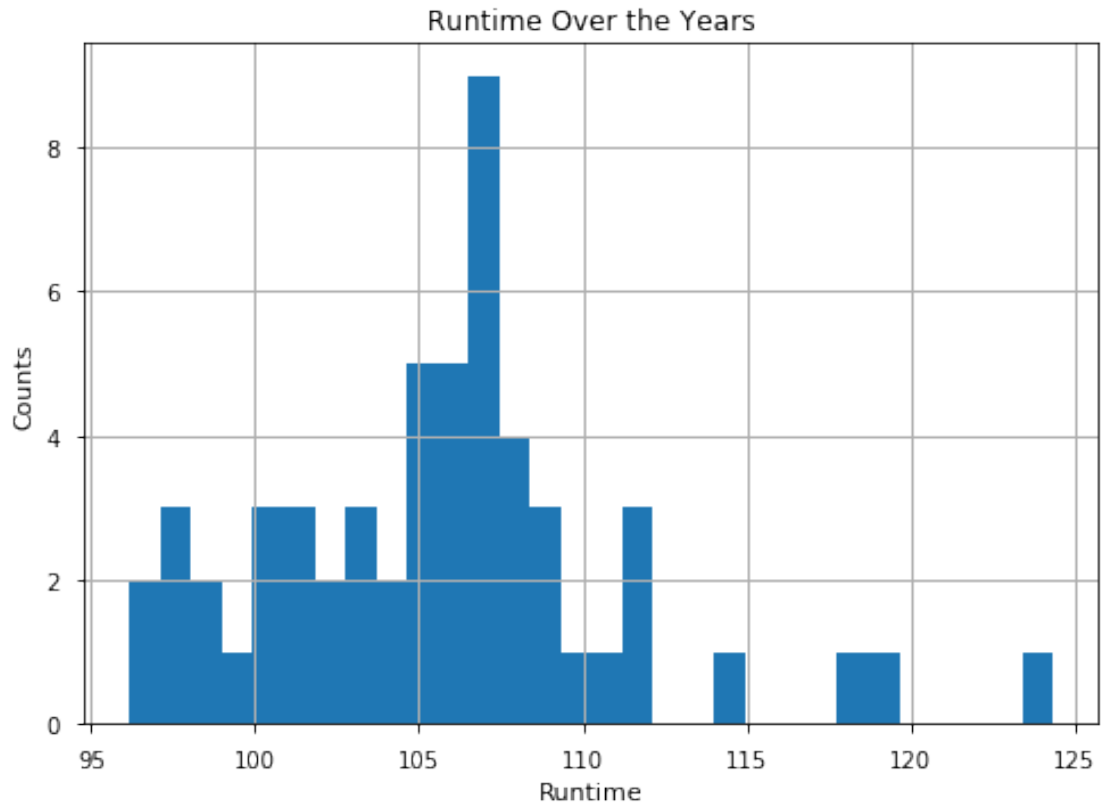
Out[84]: <seaborn.axisgrid.PairGrid at 0x1c21b6c4470>



## 5 Runtime


We wish to determine if over the years, there is popular movie running duration. Accordingly, we use a histogram to observe the counts of movie runtimes from our new dataset. Maximum counts will reflect the duration that most movies run.

```
In [69]: #exploring runtime over the years
df_new['runtime'].hist(bins=30)
plt.xlabel('Runtime')
plt.ylabel('Counts')
plt.title('Runtime Over the Years');
plt.style.use('seaborn-notebook')
plt.show()
```



In [96]: *#using seaborn distplot which helps in visualising the type of distribution*  
`sns.distplot(df_new['runtime'],color='b')`

Out[96]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1c22735bc50>



```
In [52]: # this statistic is used to answer the question mathematically
df_new['runtime'].describe()
```

```
Out[52]: count      56.000000
mean       105.714644
std        5.508458
min        96.179331
25%       101.910776
50%       105.678699
75%       107.593247
max        124.343750
Name: runtime, dtype: float64
```

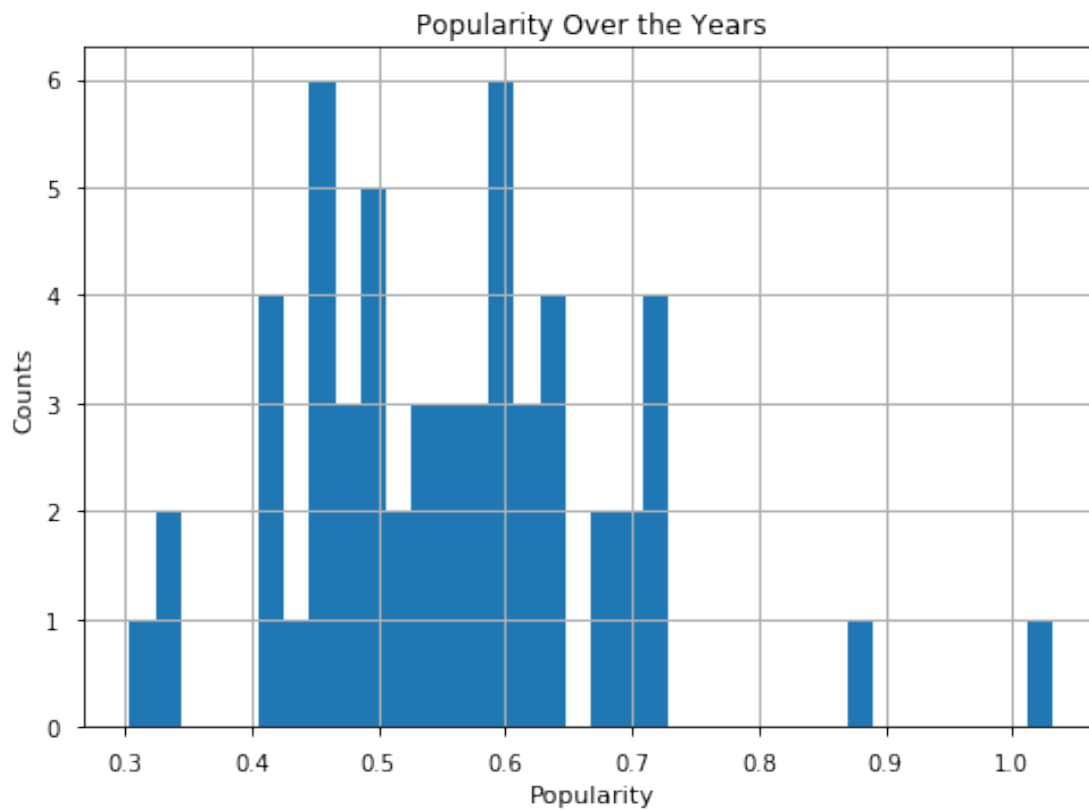
## 6 Observations about Runtimes:

As seen in the plots and functions above, popular runtimes over the years are between 106 and 107 minutes. Maximum movies fell in the 105-107 range. The distribution is right skewed.

## 7 Popularity


We wish to determine if over the years, what is the likely rating a movie will receive Accordingly, we use a histogram to observe the counts of popularity from our new dataset Maximum counts will reflect the typical popularity rating of movies

```
In [70]: ##Exploring popularity variable over the years to determine typical ratings by audien
df_new['popularity'].hist(bins=36)
plt.xlabel('Popularity')
plt.ylabel('Counts')
plt.title('Popularity Over the Years');
plt.style.use('seaborn-notebook')
plt.show()
```



```
In [98]: #using seaborn distplot which helps in visualising the type of distribution
sns.distplot(df_new['popularity'],color='g')
```

```
Out[98]: <matplotlib.axes._subplots.AxesSubplot at 0x1c2273afcf8>
```



```
In [65]: # We wish to determine if over the years, there is typical popularity rating
df_new['popularity'].describe()
```

```
Out[65]: count      56.000000
         mean        0.559693
         std         0.128434
         min         0.304112
         25%         0.469625
         50%         0.546928
         75%         0.626934
         max         1.032126
         Name: popularity, dtype: float64
```

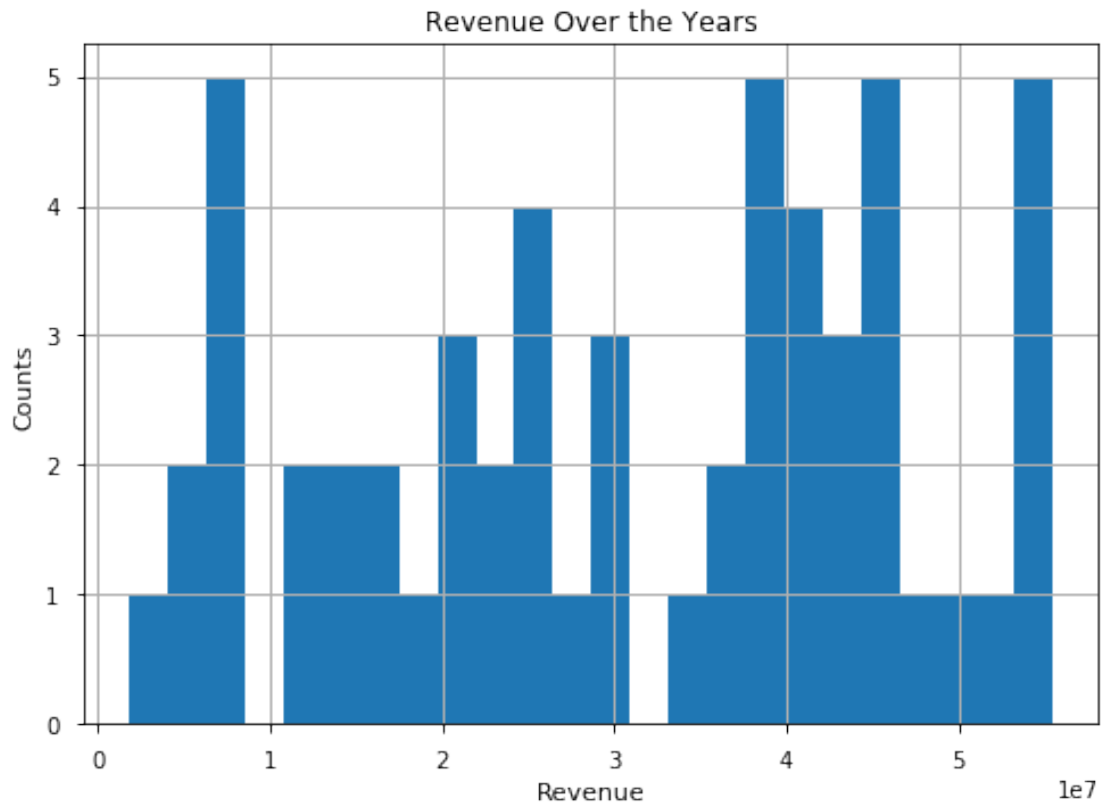
## 8 Observations about Popularity Ratings:

from histogram and quartile percentages, maximum ratings received fall in the range of 0.44 to 0.62 . The distribution is skewed to right. It is observed that no values lie in different ranges like 0.34-0.4, 0.64-0.66, 0.73-0.88, etc.Further scrutiny is required.

## 9 Revenue


We wish to determine if over the years, what is the likely revenue a movie will generate Accordingly, we use a histogram to observe the counts of revenues from our dataset Maximum counts will reflect the typical revenue earned by movies

```
In [72]: #Exploring revenue variable over the years to determine typical revenues grossed by t
df_new['revenue'].hist(bins=24)
plt.xlabel('Revenue')
plt.ylabel('Counts')
plt.style.use('seaborn-notebook')
plt.title('Revenue Over the Years');
plt.show()
```



```
In [100]: #using seaborn distplot which helps in visualising the type of distribution
sns.distplot(df_new['revenue'],color='r')
```

```
Out[100]: <matplotlib.axes._subplots.AxesSubplot at 0x1c2273e16a0>
```



```
In [73]: # We want to determine if over the years, there is typical revenue received
df_new['revenue'].describe()
```

```
Out [73]: count      5.600000e+01
          mean       3.076766e+07
          std        1.574209e+07
          min        1.842102e+06
          25%        1.815642e+07
          50%        3.257984e+07
          75%        4.293171e+07
          max        5.549569e+07
          Name: revenue, dtype: float64
```

## 10 Observations of Revenues:

From the plots and histogram, we can see that:

Distribution is left skewed. Revenues vary widely. Most movie revenues fall in the  $3.7e+07$  to  $4.8e+07$  ranges.. ‘

### 10.0.1 Research Question 2

What are variables that are associated with the revenues of movies spanning over the years? If so, by how much?



```
In [41]: #using correlation to determine factors influencing revenues
df_new.corr(method='pearson')
```

```
Out[41]:
```

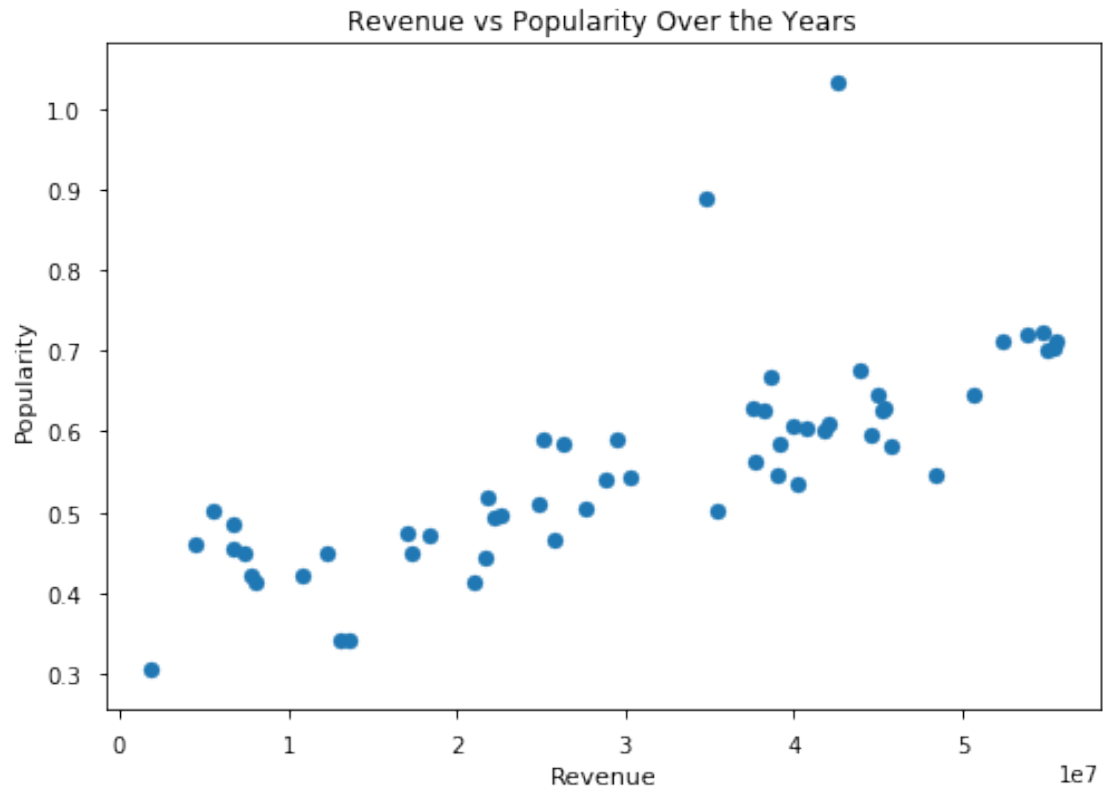
	id	popularity	budget	revenue	runtime	vote_count	\
id	1.000000	0.612610	0.100292	0.152558	-0.441775	0.539964	
popularity	0.612610	1.000000	0.668215	0.759156	-0.488974	0.849314	
budget	0.100292	0.668215	1.000000	0.906124	-0.401485	0.782026	
revenue	0.152558	0.759156	0.906124	1.000000	-0.466239	0.809243	
runtime	-0.441775	-0.488974	-0.401485	-0.466239	1.000000	-0.612715	
vote_count	0.539964	0.849314	0.782026	0.809243	-0.612715	1.000000	
vote_average	-0.297719	-0.574979	-0.731797	-0.706442	0.524770	-0.658819	
budget_adj	-0.148336	0.458952	0.891925	0.767164	-0.073168	0.505434	
revenue_adj	-0.299723	-0.100506	-0.158415	0.074952	0.300077	-0.215568	

	vote_average	budget_adj	revenue_adj
id	-0.297719	-0.148336	-0.299723
popularity	-0.574979	0.458952	-0.100506
budget	-0.731797	0.891925	-0.158415
revenue	-0.706442	0.767164	0.074952
runtime	0.524770	-0.073168	0.300077
vote_count	-0.658819	0.505434	-0.215568
vote_average	1.000000	-0.557569	0.377204
budget_adj	-0.557569	1.000000	0.050086
revenue_adj	0.377204	0.050086	1.000000


## 11 Revenue vs Popularity:

```
In [76]: # Creating a scatterplot of revenue and popularity
plt.scatter(x=df_new['revenue'], y=df_new['popularity'])
plt.xlabel('Revenue')
plt.ylabel('Popularity')
plt.title('Revenue vs Popularity Over the Years');
```



```
In [114]: #creating a join plot using seaborn
sns.jointplot(x=df_new['revenue'], y=df_new['popularity'],color="orange")
```

```
Out[114]: <seaborn.axisgrid.JointGrid at 0x1c228bc8588>
```




## 12 Observations- Revenue vs Popularity:

The correlation computed from function is evident in the scatterplot. Popularity is proportional to revenues. The few outliers that grossed medium but were rated extremely high need to be looked into.

## 13 Revenue vs Vote counts:


```
In [102]: # Creating a scatterplot of revenue and vote counts over the years
plt.scatter(x=df_new['revenue'], y=df_new['vote_count'])
plt.xlabel('Revenue')
```

```
plt.ylabel('Vote Count')
plt.title('Revenue vs Vote Count Over the Years');
```



```
In [115]: #creating a join plot using seaborn
sns.jointplot(x=df_new['revenue'], y=df_new['vote_count'],color="blue")
```

```
Out[115]: <seaborn.axisgrid.JointGrid at 0x1c228a6eda0>
```



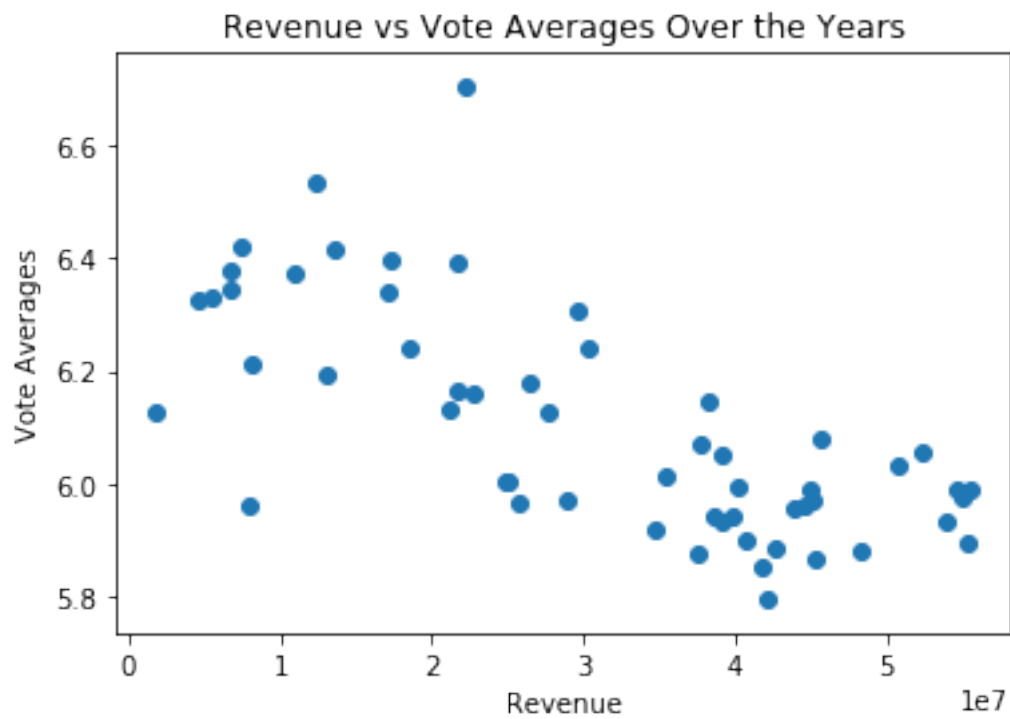
## 14 Observations- Revenue vs Vote Count:

The correlation computed from function is evident in the scatterplot. Vote count is roughly proportional to revenues. Even in this plot, we observe that a few sample had high counts but did not gross commensurately in revenue. This plot points us to scrutinize the vote counts and type of votes if possible or any other factors influencing voting such as genres, cast, etc.

## 15 Revenue vs Vote averages:


```
In [47]: # Creating a scatterplot of revenue and vote over the years
plt.scatter(x=df_new['revenue'], y=df_new['vote_average'])
plt.xlabel('Revenue')
```

```
plt.ylabel('Vote Averages')
plt.title('Revenue vs Vote Averages Over the Years');
```



```
In [117]: #creating a join plot using seaborn
sns.jointplot(x=df_new['revenue'], y=df_new['vote_average'],color="green")
```

```
Out[117]: <seaborn.axisgrid.JointGrid at 0x1c229f53668>
```



## 16 Observations- Revenue vs Vote Average:

The correlation computed from function is evident in the scatterplot. Vote averages are inversely proportional to revenues. Considering this plot and the previous, we can infer voting variables as influencing revenues. This means that detailed analysis is required to assess the outliers and possibly pursue a detailed line of investigation based on vote counts and averages.

## 17 Limitations:

1. In above analysis, we have seen runtimes, popularity and revenues.

2. We have restricted ourselves to only numerical values of revenues. Any variances in currencies, like dollars, rupees, Euros, etc. have not been considered.

3. In the case of varying currencies which is likely as the movies ranges from different regions, exchange rates or conversions or some sort of normalization needed to be applied.

3. We have not considered different revenue sources like movies audio rights sales, ticket sales, digital sales, broadcast sales, etc. and disregard the period and markets over which the revenues are collected. Older movies that are casted over television or re-released which might garnered higher revenues are examples of limitations of the dataset and hence, this EDA.

4. Vote counts have been taken at face value. We are not inspecting the nature of votes, good or bad or any other range of values. Also, the period or channels in which these vote counts were gathered is not considered and will be significant if varying for different movies.

5. The medium through which popularity is determined is unknown. This can impact the analysis as the limitations and bias inherent while gauging audience response will be present in the end values as well.

6. This analysis assumes that the same methods and index were employed for collecting popularity factors and counting votes for all the movies. In the event that it is not so, the results might not hold true. A possibility since movies are from different countries and languages.

7. While we did not have missing values for any of the factors taken under consideration, we can acknowledge the presence of these limitations and assumptions in our analysis of the dataset.

#### ## Conclusions

1. From our analysis, we discovered that over the years, there are patterns to runtimes, popularity and revenues. While only tentative, we have found that popular runtimes range between 90 and 100 minutes.

2. Revenues showed a wider range, but the most likely range was  $8.838806 \times 10^7$  to  $1.028258 \times 10^8$ .

3. We also analyzed multiple variables. Specifically, we chose to analyze over time [release\_date], the effect of particular variables [votes and popularity] on our factor of interest [revenue].

4. Based on this analysis, we found that over the years, popularity and revenue show a direct relation. Vote counts and revenue are related positively while vote averages are inversely related to revenues. However, these relations are merely correlations and do not imply causation.

5. These lines of analysis point to the need for further investigation, especially with votes.

## 18 Thank you